

Analyzing the Internet Using DNS

Initial Bachelor Talk

Lennart Bader

Advisor: Dr. Oliver Hohlfeld

<http://comsys.rwth-aachen.de>

Aachen, 2016-12-12

Motivation

Domain name system



Motivation

Domain name system

- Central part of the Internet



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI
- Internet evolution: IPv6 deployment



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI
- Internet evolution: IPv6 deployment
- Security, Validity, ...



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI
- Internet evolution: IPv6 deployment
- Security, Validity, ...
- Some aspects already in current research:



Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI
- Internet evolution: IPv6 deployment
- Security, Validity, ...
- Some aspects already in current research:
DNSSEC Analysis [IMC'14]



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI
- Internet evolution: IPv6 deployment
- Security, Validity, ...
- Some aspects already in current research:
DNSSEC Analysis [IMC'14], Webserver Security [IMC'14]



Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI
- Internet evolution: IPv6 deployment
- Security, Validity, ...
- Some aspects already in current research:
DNSSEC Analysis [IMC'14], Webserver Security [IMC'14],
DNS Validity [IMC'15], ...



Motivation

Domain name system

- Central part of the Internet
- Allows analyses on several aspects

Possibilities

- DNS "Junk" Analysis
- Support for ZMap Port Scans to enable SNI
- Internet evolution: IPv6 deployment
- Security, Validity, ...
- Some aspects already in current research:
DNSSec Analysis [IMC'14], Webserver Security [IMC'14],
DNS Validity [IMC'15], ...



Goal: Collect DNS data periodically for many domains

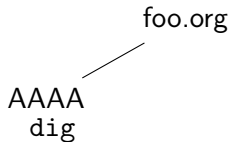
Seems easy...

foo.org

dig

Seems easy...

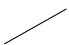
foo.org
AAAA
dig



Just needed $\approx 36ms$

Seems easy...

foo.org
AAAA
dig



Just needed $\approx 36ms$

Where is the challenge?


Seems easy...




Seems easy...




Seems easy...



Seems easy...




Seems easy...



Internet is huge


Seems easy...



Internet is huge

- Millions of domains
- Different analyses (not only IPv6)

Seems easy...




Internet is huge

- Millions of domains
- Different analyses (not only IPv6)


⇒ Multiple Queries per domain + Millions of domains

Seems easy...



$$N_{\text{Queries}} \geq (\underbrace{130M}_{\text{.com}} + \underbrace{15M}_{\text{.net}} + \underbrace{10M}_{\text{.org}} + \underbrace{\dots}_{\text{Other TLDs}}) * N_{\text{Queries per Domain}}$$

Seems easy...



$$N_{\text{Queries}} \geq (\underbrace{130M}_{\text{.com}} + \underbrace{15M}_{\text{.net}} + \underbrace{10M}_{\text{.org}} + \underbrace{\dots}_{\text{Other TLDs}}) * N_{\text{Queries per Domain}}$$

$$\Rightarrow N_{\text{Queries}} \geq 1.55 \text{ Billion}$$

Seems easy...

```
dig www.rwth-aachen.de
```

Seems easy...

```
dig www.rwth-aachen.de
```

- `www.rwth-aachen.de. 99054 IN A 137.226.107.63`

Seems easy...

```
dig www.rwth-aachen.de
```

- `www.rwth-aachen.de. 99054 IN A 137.226.107.63`
- Just needed $\approx 36ms$

Seems easy...

```
dig www.rwth-aachen.de
```

- `www.rwth-aachen.de. 99054 IN A 137.226.107.63`
- Just needed $\approx 36ms$

```
dig www.rwth-aachen.de AAAA
```

Seems easy...

```
dig www.rwth-aachen.de
```

- `www.rwth-aachen.de. 99054 IN A 137.226.107.63`
- Just needed $\approx 36ms$

```
dig www.rwth-aachen.de AAAA
```

- RWTH has no AAAA record...

Seems easy...

```
dig www.rwth-aachen.de
```

- `www.rwth-aachen.de. 99054 IN A 137.226.107.63`
- Just needed $\approx 36ms$

```
dig www.rwth-aachen.de AAAA
```

- RWTH has no AAAA record...
- Again $\approx 36ms$

Seems easy...

```
dig www.rwth-aachen.de
```

- `www.rwth-aachen.de. 99054 IN A 137.226.107.63`
- Just needed $\approx 36ms$

```
dig www.rwth-aachen.de AAAA
```

- RWTH has no AAAA record...
- Again $\approx 36ms$

Where is the challenge?

Analyzing one domain

E.g.: Get information about IPv6 compatibility


Analyzing one domain

E.g.: Get information about IPv6 compatibility

foo.org

Analyzing one domain

E.g.: Get information about IPv6 compatibility



Analyzing one domain

E.g.: Get information about IPv6 compatibility




Analyzing one domain

E.g.: Get information about IPv6 compatibility




Analyzing one domain

E.g.: Get information about IPv6 compatibility




Analyzing one domain

E.g.: Get information about IPv6 compatibility



Analyzing one domain

E.g.: Get information about IPv6 compatibility



⇒ Multiple Queries per Domain

Analyzing the Internet

Internet is huge

- Millions of domains
- Different analyses (not only IPv6)

Internet is huge

- Millions of domains
- Different analyses (not only IPv6)

Different TLDs

- .com, .org, .net, ...
- Millions of domains per TLD
- Multiple queries per domain

Analyzing the Internet

Internet is huge

- Millions of domains
- Different analyses (not only IPv6)

Different TLDs

- .com, .org, .net, ...
- Millions of domains per TLD
- Multiple queries per domain

$$N_{\text{Queries}} \geq (\underbrace{130M}_{\text{.com}} + \underbrace{15M}_{\text{.net}} + \underbrace{10M}_{\text{.org}} + \underbrace{\dots}_{\text{Other TLDs}}) * N_{\text{Queries per Domain}}$$

Analyzing the Internet

Internet is huge

- Millions of domains
- Different analyses (not only IPv6)

Different TLDs

- .com, .org, .net, ...
- Millions of domains per TLD
- Multiple queries per domain

$$N_{\text{Queries}} \geq (\underbrace{130M}_{\text{.com}} + \underbrace{15M}_{\text{.net}} + \underbrace{10M}_{\text{.org}} + \underbrace{\dots}_{\text{Other TLDs}}) * N_{\text{Queries per Domain}}$$

$$\Rightarrow N_{\text{Queries}} \geq 1.55 \text{ Billion}$$

dig does not scale

- 100.000 queries take about 4 hours
- \Rightarrow Querying 130M domains: ≈ 216 days

dig does not scale

- 100.000 queries take about 4 hours
- \Rightarrow Querying 130M domains: ≈ 216 days

Scalability is a big challenge!

dig does not scale

- 100.000 queries take about 4 hours
- \Rightarrow Querying 130M domains: \approx 216 days

Scalability is a big challenge!

Can we do better?

Enable performant large scale analyses on the Internet




Enable performant large scale analyses on the Internet

- Large scale:
Hundreds of Millions of domains



Enable performant large scale analyses on the Internet

- Large scale:
Hundreds of Millions of domains
- Multiple Analyses:
IPv6 deployment, Port Scans with SNI, ...




Enable performant large scale analyses on the Internet

- Large scale:
Hundreds of Millions of domains
- Multiple Analyses:
IPv6 deployment, Port Scans with SNI, ...
- Long term evolution:
Repeat analyses periodically over years



Enable performant large scale analyses on the Internet

- Large scale:
Hundreds of Millions of domains
- Multiple Analyses:
IPv6 deployment, Port Scans with SNI, ...
- Long term evolution:
Repeat analyses periodically over years
- Performance:
Be much faster than dig



- Enable performant DNS resolution in large scale

- Enable performant DNS resolution in large scale
 - Investigate performance:
Time usage

- Enable performant DNS resolution in large scale
 - Investigate performance:
Time usage
 - Investigate suitability:
Reliability / success rate

- Enable performant DNS resolution in large scale
 - Investigate performance:
Time usage
 - Investigate suitability:
Reliability / success rate
- Analysis Framework:
 - Enable long term and large scale analyses

- Enable performant DNS resolution in large scale
 - Investigate performance:
Time usage
 - Investigate suitability:
Reliability / success rate
- Analysis Framework:
 - Enable long term and large scale analyses
 - Analysis: Investigate specific aspects of the Internet

- Enable performant DNS resolution in large scale
 - Investigate performance:
Time usage
 - Investigate suitability:
Reliability / success rate
- Analysis Framework:
 - Enable long term and large scale analyses
 - Analysis: Investigate specific aspects of the Internet
- Analyze current IPv6 deployment

- Enable performant DNS resolution in large scale
 - Investigate performance:
Time usage
 - Investigate suitability:
Reliability / success rate
- Analysis Framework:
 - Enable long term and large scale analyses
 - Analysis: Investigate specific aspects of the Internet
- Analyze current IPv6 deployment
- (Optional) analyze further aspects (e.g. Junk Analysis)

First Results

Tested performance and suitability of ZDNS
(released 5 months ago)



First Results

Tested performance and suitability of ZDNS
(released 5 months ago)

- Queried 1 Million domains (Alexa)



First Results

Tested performance and suitability of ZDNS
(released 5 months ago)

- Queried 1 Million domains (Alexa)
- Measured needed time (performance)



First Results

Tested performance and suitability of ZDNS
(released 5 months ago)

- Queried 1 Million domains (Alexa)
- Measured needed time (performance)
- Investigated success rate



First Results

Tested performance and suitability of ZDNS
(released 5 months ago)

- Queried 1 Million domains (Alexa)
- Measured needed time (performance)
- Investigated success rate
- Evaluated result quality:



First Results

Tested performance and suitability of ZDNS
(released 5 months ago)

- Queried 1 Million domains (Alexa)
- Measured needed time (performance)
- Investigated success rate
- Evaluated result quality:
 - Completeness of results



First Results

Tested performance and suitability of ZDNS
(released 5 months ago)

- Queried 1 Million domains (Alexa)
- Measured needed time (performance)
- Investigated success rate
- Evaluated result quality:
 - Completeness of results
 - Intermediate steps of resolution?



First Results

Tested performance and suitability of ZDNS
(released 5 months ago)

- Queried 1 Million domains (Alexa)
- Measured needed time (performance)
- Investigated success rate
- Evaluated result quality:
 - Completeness of results
 - Intermediate steps of resolution?



Compared results with those obtained from own tool.

First Results

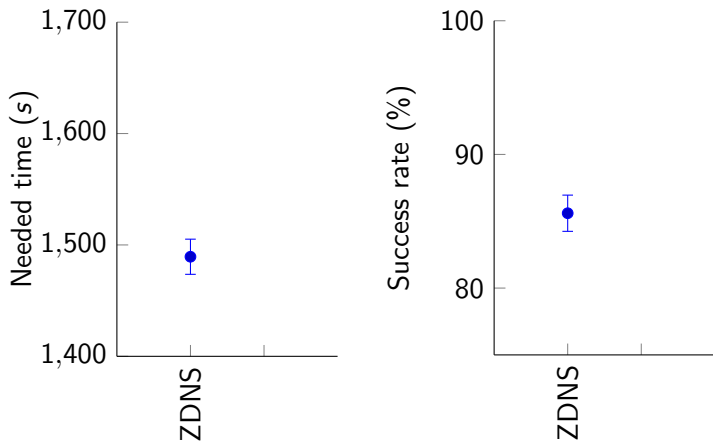


Figure: Test results for ZDNS
(1.000.000 queries)

First Results




Figure: Test results for ZDNS and TDNS
(1.000.000 queries)

Timeline

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16

Timeline

Implementation

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16


Timeline

Tool Analyses


Implementation

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16


Timeline




Timeline



Timeline



Timeline



References



John S. Otto, Mario A. Sánchez, John P. Rula, and Fabián E. Bustamante.

Content delivery and the natural evolution of dns.



Tom van Goethem, Ping Chen, Nick Nikiforakis, Lieven Desmet, and Wouter Joosen.

Large-scale security analysis of the web: Challenges and findings.



Roland van Rijswijk-Deij, Mattijs Jonker, Anna Sperotto, and Aiko Pras.

A high-performance, scalable infrastructure for large-scale active dns measurements.



Roland van Rijswijk-Deij, Anna Sperotto, and Aiko Pras.
Dnssec and its potential for ddos attacks.